

```
!pip install pandas yfinance scikit-learn matplotlib plotly
```

```

Requirement already satisfied: pandas in /usr/local/lib/python3.10/dist-packages (2.2.2)
Requirement already satisfied: yfinance in /usr/local/lib/python3.10/dist-packages (0.2.50)
Requirement already satisfied: scikit-learn in /usr/local/lib/python3.10/dist-packages (1.5.2)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.10/dist-packages (3.8.0)
Requirement already satisfied: plotly in /usr/local/lib/python3.10/dist-packages (5.24.1)
Requirement already satisfied: numpy>=1.22.4 in /usr/local/lib/python3.10/dist-packages (from pandas)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.10/dist-packages (from pandas)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.10/dist-packages (from pandas)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.10/dist-packages (from pandas)
Requirement already satisfied: requests>=2.31 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: multitasking>=0.0.7 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: lxml>=4.9.1 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: platformdirs>=2.0.0 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: frozendict>=2.3.4 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: peewee>=3.16.2 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: beautifulsoup4>=4.11.1 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: html5lib>=1.1 in /usr/local/lib/python3.10/dist-packages (from yfinance)
Requirement already satisfied: scipy>=1.6.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn)
Requirement already satisfied: joblib>=1.2.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn)
Requirement already satisfied: threadpoolctl>=3.1.0 in /usr/local/lib/python3.10/dist-packages (from scikit-learn)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: cycler>=0.10 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: kiwisolver>=1.0.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: pillow>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.10/dist-packages (from matplotlib)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.10/dist-packages (from plotly)
Requirement already satisfied: soupsieve>1.2 in /usr/local/lib/python3.10/dist-packages (from beautifulsoup4)
Requirement already satisfied: six>=1.9 in /usr/local/lib/python3.10/dist-packages (from html5lib)
Requirement already satisfied: webencodings in /usr/local/lib/python3.10/dist-packages (from html5lib)
Requirement already satisfied: charset-normalizer<4,>=2 in /usr/local/lib/python3.10/dist-packages (from requests)
Requirement already satisfied: idna<4,>=2.5 in /usr/local/lib/python3.10/dist-packages (from requests)
Requirement already satisfied: urllib3<3,>=1.21.1 in /usr/local/lib/python3.10/dist-packages (from requests)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.10/dist-packages (from requests)

```

```

import yfinance as yf
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.ensemble import RandomForestRegressor
from sklearn.metrics import mean_squared_error, mean_absolute_error
from sklearn.preprocessing import StandardScaler, minmax_scale

#define a function to fetch stock data

def fetch_stock_data(stock_ticker, start_date="2020-01-01", end_date="2024-01-01"):
    try:
        return yf.download(stock_ticker, start=start_date, end=end_date).reset_index()
    except Exception as e:
        print(f"Error: {e}")
        return None

# Example usage
df = fetch_stock_data("TSLA")

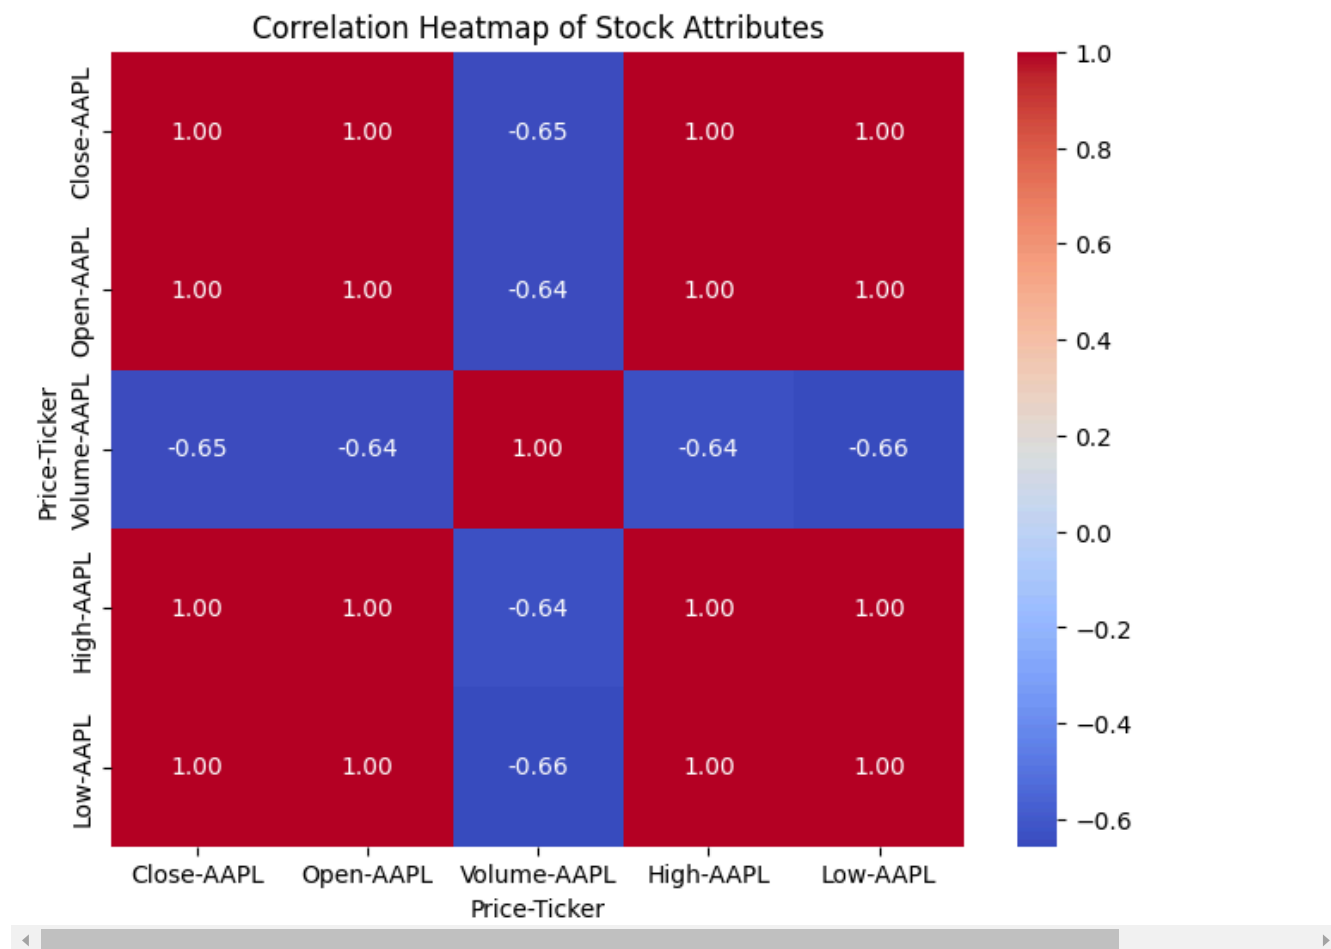
```

df

```
def plot_correlation_heatmap(data):  
    selected_columns = ['Close', 'Open', 'Volume', 'High', 'Low']  
    corr = data[selected_columns].corr() # Calculate correlation matrix  
  
    # Display correlation values as a heatmap  
    plt.figure(figsize=(8, 6))  
    sns.heatmap(corr, annot=True, cmap='coolwarm', fmt=".2f")  
    plt.title("Correlation Heatmap of Stock Attributes")  
    plt.show()
```

```
stock_data = fetch_stock_data("AAPL")  
plot_correlation_heatmap(stock_data)
```

➡ [*****100%*****] 1 of 1 completed



Strong Positive Correlation (Red, Value = 1.00):

Close-AAPL, Open-AAPL, High-AAPL, and Low-AAPL all have a perfect positive correlation (1.00) with each other. This means these price-related attributes move in the same direction. If one increases, the others also tend to increase proportionally.


Reason: Since "Close," "Open," "High," and "Low" are directly related to stock prices, their high correlation is natural. They are part of the same data series for a stock.

Moderate Negative Correlation (Blue, ~ -0.64 to -0.66):

Volume-AAPL has a moderate negative correlation with all price attributes: Volume-AAPL vs Close-AAPL: -0.65
Volume-AAPL vs Open-AAPL: -0.64 Volume-AAPL vs High-AAPL: -0.64 Volume-AAPL vs Low-AAPL: -0.66


Interpretation: This indicates that when Volume (number of shares traded) increases, the price-related metrics (Open, Close, High, Low) tend to decrease slightly. This might suggest: Increased selling activity during high-volume days, which pushes prices down. Volatility in the market.

```
plt.figure(figsize=(10, 6))
plt.plot(stock_data['Date'], stock_data['Close'], label='Close Price')
plt.title('Stock Close Price Over Time')
plt.xlabel('Date')
plt.ylabel('Close Price')
plt.legend()
plt.show()
```



```
# Calculate volatility (30-day rolling standard deviation)
stock_data['Volatility'] = stock_data['Close'].rolling(window=30).std()

# Plot volatility over time
plt.figure(figsize=(10, 6))
plt.plot(stock_data['Date'], stock_data['Volatility'], color='orange', label='Volatility')
plt.title('Volatility of Stock Close Price')
plt.xlabel('Date')
plt.ylabel('Volatility')
plt.legend()
plt.show()
```



Measuring volatility of the stock

High Peaks: Peaks (e.g., mid-2020, early 2022, mid-2022) indicate high volatility. High volatility often occurs during major events like market crashes, earnings reports, economic uncertainty, or external shocks (e.g., COVID-19 impact in 2020). **Low Points:** Periods where the volatility drops indicate more stable price movements. This could mean reduced uncertainty and consistent trading behavior. **Pattern:** The chart shows recurring cycles of volatility, suggesting that stock price fluctuations often follow events or trends. **Practical Implications:**

**High volatility ** → Higher risk and potential for large price swings. Low volatility → More stable price trends, making the stock less risky.

**1. Predict Future Closing Price ** Train a model to predict the next day's closing price based on historical features like Open, High, Low, Volume, and previous Close prices.

```
def prepare_data(stock_data):
    # Add "Previous Close" feature
    stock_data['Prev_Close'] = stock_data['Close'].shift(1)

    # Remove NaN from the first row
    stock_data = stock_data.dropna()

    # Features and target
    X = stock_data[['Prev_Close', 'Open', 'High', 'Low', 'Volume']]
    y = stock_data['Close'] # Target: next day's Close price

    return X, y, stock_data['Date']
```

```
from sklearn.metrics import r2_score
def Linear_Regression_check(X, y, dates):
    # Split the data into training and test sets
```

```
# Split the data into train and test sets
X_train, X_test, y_train, y_test, dates_train, dates_test = train_test_split(
    X, y, dates, test_size=0.2, random_state=42, shuffle=False
)

# Initialize and train the model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on the test set
y_pred = model.predict(X_test)

# Evaluate the model
mae = mean_absolute_error(y_test, y_pred)
r2 = r2_score(y_test, y_pred)
print(f"Mean Absolute Error: {mae:.2f}")
print(f"R-squared: {r2:.2f}")


# Plot actual vs predicted prices
plt.figure(figsize=(10, 6))
plt.plot(dates_test, y_test, label='Actual Price', color='blue')
plt.plot(dates_test, y_pred, label='Predicted Price', color='red')
plt.title("Actual vs Predicted Closing Prices")
plt.xlabel("Date")
plt.ylabel("Closing Price")
plt.legend()
plt.show()

if __name__ == "__main__":
    # Step 1: Fetch the stock data
    ticker = "AAPL" # Example stock ticker: Apple Inc.
    stock_data = fetch_stock_data(ticker)

    # Step 2: Prepare the data
    X, y, dates = prepare_data(stock_data)

    # Step 3: Train and evaluate the model
    LinearRegression_check(X, y, dates)
```

➡ [*****100%*****] 1 of 1 completed
 Mean Absolute Error: 0.62
 R-squared: 0.99



Mean Absolute Error (MAE): 0.62

This is an excellent result since the error is very low. On average, the model's prediction deviates by only 0.62 from the actual closing prices which is minimal for stock prices in the 150-\$200 range.

**R-squared* (R^2): *0.99 An R^2 value of 0.99 indicates the model explains 99% of the variance in the data. This is an extremely strong fit, meaning the model is performing very well.

SUCH AN ACCURATE MAE AND R2 ERROR COULD SHOW OVERFITTING , LETS CHECK FOR THAT

```
from sklearn.model_selection import TimeSeriesSplit, cross_val_score
```

```
tscv = TimeSeriesSplit(n_splits=5)
model = LinearRegression()
scores = cross_val_score(model, X, y, cv=tscv, scoring='r2')
print("Cross-validation R2 scores:", scores)
print("Mean R2 score:", scores.mean())
```

➡ Cross-validation R2 scores: [0.9838707 0.99576748 0.9909305 0.99046169 0.99171093]
 Mean R2 score: 0.9905482585553159

The Cross-Validation R^2 scores and the mean R^2 score (0.99) indicate that your model is performing consistently across different folds of the data. This reduces the likelihood of overfitting.

2. Predict Stock Volatility Goal: Predict the stock's future volatility (e.g., 30-day rolling standard deviation) based on past price movements and volume.

```
import xgboost as xgb

def prepare_volatility_data(stock_data):
    # Calculate 30-day rolling standard deviation as the target (Volatility)
    stock_data['Volatility'] = stock_data['Close'].rolling(window=30).std()

    # Add lagged features
    stock_data['Volatility_Lag1'] = stock_data['Volatility'].shift(1)

    # Add rolling statistics for Volume
    stock_data['Volume_SMA_10'] = stock_data['Volume'].rolling(window=10).mean()
    stock_data['Volume_SMA_30'] = stock_data['Volume'].rolling(window=30).mean()

    # Add SMA for prices
    stock_data['SMA_10'] = stock_data['Close'].rolling(window=10).mean()
    stock_data['SMA_20'] = stock_data['Close'].rolling(window=20).mean()

    # Drop rows with NaN
    stock_data = stock_data.dropna()

    # Features and target
    X = stock_data[['Prev_Close', 'Open', 'High', 'Low', 'Volume',
                    'SMA_10', 'SMA_20', 'Volume_SMA_10', 'Volume_SMA_30', 'Volatility_Lag1']]
    y = stock_data['Volatility']
    return X, y, stock_data['Date']

# Train and evaluate XGBoost model
def train_and_evaluate_xgboost(X, y, dates):
    # Sequential Train-Test Split
    X_train, X_test, y_train, y_test, dates_train, dates_test = train_test_split(
        X, y, dates, test_size=0.2, random_state=42, shuffle=False
    )

    # Train XGBoost Regressor
    model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1, random_state=42)
    model.fit(X_train, y_train)

    # Predictions
    y_pred = model.predict(X_test)


    # Evaluation
    mae = mean_absolute_error(y_test, y_pred)
    r2 = r2_score(y_test, y_pred)

    print(f"Mean Absolute Error: {mae:.4f}")
    print(f"R-squared: {r2:.4f}")

    # Plot Actual vs Predicted Volatility
    plt.figure(figsize=(10, 6))
    plt.plot(dates_test, y_test, label='Actual Volatility', color='blue')
    plt.plot(dates_test, y_pred, label='Predicted Volatility', color='red')
    plt.title("Actual vs Predicted Volatility using XGBoost")
    plt.xlabel("Date")
    plt.ylabel("Volatility")
    plt.legend()
    plt.show()
```

```
if __name__ == "__main__":  
    ticker = "AAPL" # Example stock  
    stock_data = fetch_stock_data(ticker)  
  
    # Add previous day's close for consistency with earlier features  
    stock_data['Prev_Close'] = stock_data['Close'].shift(1)  
  
    # Prepare features and target  
    X, y, dates = prepare_volatility_data(stock_data)  
  
    # Train and evaluate the model  
    train_and_evaluate_xgboost(X, y, dates)
```

➡ [*****100%*****] 1 of 1 completed
Mean Absolute Error: 0.3124
R-squared: 0.9311



stock_data



Price	Date	Adj Close	Close	High	Low	Open	Volume	Prev_Close	Volat
Ticker		AAPL	AAPL	AAPL	AAPL	AAPL	AAPL		
0	2020-	72.796013	75.087502	75.150002	73.797501	74.059998	135480400	NaN	